

DoH-capable DNS forwarder

This project will require using Python to build a simple DNS forwarder with domain blocking and DoH capabilities. This DNS forwarder will need to do the following: (1) receive an arbitrary DNS message from a client, (2) check if the domain name should be blocked, and if so respond with an NXDomain message, (3) if the queried domain name is allowed, forward the DNS message to either standard DNS resolver or a DoH-capable resolver, (4) wait for the response from the resolver and forward it back to the client.

=== Details ===

Schema:

```
Client <==> DNS forwarder <==> DNS resolver < = = =  
> (DNS NSes)  
          ^^^^^^^^^^^^^^^
```

Command line parameters:

```
$ ./dns_forwarder.py -h  
usage: dns_forwarder.py [-h] [-d DST_IP] -f  
DENY_LIST_FILE  
                        [-l LOG_FILE] [--doh] [--  
doh_server DOH_SERVER]  
optional arguments:  
-h, --help      show this help message and exit  
-d DST_IP       Destination DNS server IP  
-f DENY_LIST_FILE File containing domains to block  
-l LOG_FILE      Append-only log file  
--doh           Use default upstream DoH server  
--doh_server DOH_SERVER Use this upstream DoH  
server
```

Requirements:

- If `--doh` or `--doh_server` are specified, the forwarder **MUST** forward the DNS query using the DoH protocol
- If `--doh` or `--doh_server` are not specified (in which case `-d` **MUST** be present), the forwarder **MUST** forward the DNS query using the DNS protocol

- The DNS forwarder MUST receive DNS messages from the client via a simple UDP server socket.
- When DoH is not used, the `-d` option will be specified and the forwarder must use a simple UDP client socket to forward the client's query to the DNS resolver
- The `DENY_LIST_FILE` file MUST contain a (potentially empty) list of domain names that MUST be blocked by the forwarder.

DoH REQUESTS

You are required to use GET requests as defined in RFC 8484. For instance, a DNS request for `(qname='example.com', qtype='A')` would look like:

`https://8.8.8.8/dns-query?dns=AAABAAABAAAAAAB2V4YW1wbGUDY29tAAABAAE`

Note: do not use the JSON API provided by some DNS operators. You can find more information here:

- <https://developers.google.com/speed/public-dns/docs/secure-transport#doh>
- <https://developers.cloudflare.com/1.1.1.1/encrypted-dns/dns-over-https/make-api-requests/dns-wireformat>
- <https://datatracker.ietf.org/doc/html/rfc8484>

NOTE: when working from the UGA campus network (including from your VM), use IP addresses (e.g., 1.1.1.1, 8.8.8.8, etc.) in the DoH query URL, because it appears that EITS is blocking some domain name strings related to DoH servers (e.g., dns.google or cloudflare-dns.com).

DENY LIST FORMAT

The deny list provided in input will be a text file (not necessarily ending in .txt) containing one domain name per line. For instance:

```
www.example.com
cobweb.cs.uga.edu
yahoo.co.jp
```

only fully qualified domains should be blocked. For instance, in the above example only a domain matching yahoo.co.jp must be blocked (e.g., www.yahoo.co.jp should not be blocked).

Note: when a domain is blocked, your DNS forwarder MUST reply to the client with a properly formatted NXDOMAIN response.

LOG FILE ENTRY FORMAT

The log file should be a text file containing a record of all domain names and query types that have been requested, and whether the request was blocked or allowed. For instance:

```
www.google.com A ALLOW
google.com NS ALLOW
www.yahoo.co.jp A DENY
yahoo.co.jp MX DENY
www.youtube.com A ALLOW
www.example.com A DENY
...
```

SUBMISSION GUIDELINES

Create a directory named with your last name. Put your program's files under that directory. Then, create a Project3-FIRSTNAME-LASTNAME.tar.gz archive containing that directory. For instance, in my case those would be:

```
_ perdisci
|___ dns_forwarder.py
```

Project3-ROBERTO-PERDISCI.tar.gz

Finally, submit the .tar.gz file through eLC.

Here is an example of how the code will be executed:

```
$ ./dns_forwarder.py --doh_server 1.1.1.1 -l
./queries.log -f ./deny_list.txt
```

Assuming dns_fowarder.py is running on vmX.cs.uga.edu, you could use dig to test it. For instance (this is only an example):

```
$ dig @vmX.cs.uga.edu -t AAAA www.example.com
```

and you should see a response similar to this one (again, this is only an example):

```
; <<>> DiG 9.10.6 <<>> @vmX.cs.uga.edu -t
AAAA www.example.com
; (1 server found)
;; global options: +cmd
```

```
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id:
49313
;; flags: qr rd ra ad; QUERY: 1, ANSWER: 1,
AUTHORITY: 0, ADDITIONAL: 1
;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags;; udp: 1232
;; QUESTION SECTION:
;www.example.com.          IN      AAAA
;; ANSWER SECTION:
www.example.com.          82730    IN      AAAA      2606:2800
:220:1:248:1893:25c8:1946
;; Query time: 10 msec
;; SERVER: vmX.cs.uga.edu#53(172.17.152.X)
;; WHEN: Thu Oct 14 16:01:17 EDT 2021
;; MSG SIZE rcvd: 72
```

If www.example.com needs to be blocked (i.e., it's in the deny list), then the DNS forwarder should generate a response that looks like this

(note: this is only an example):

```
; <<>> DiG 9.10.6 <<>> @vmX.cs.uga.edu -t
AAAA www.example.com
; (1 server found)
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NXDOMAIN, id:
20003
;; flags: qr rd ra; QUERY: 1, ANSWER: 0, AUTHORITY:
0, ADDITIONAL: 1
;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags;; udp: 1232
;; QUESTION SECTION:
;www.example.com.          IN      AAAA
;; Query time: 48 msec
;; SERVER: vmX.cs.uga.edu#53(172.17.152.X)
;; WHEN: Thu Oct 14 16:08:26 EDT 2021
;; MSG SIZE rcvd: 47
```

Given that a few students are struggling with some parts of Project 3, here are some tips.

1) Basic examples on creating and parsing DNS messages with Scapy:

Example 1:

```
dns_req = DNS(rd=1, qd=DNSQR(qname='www.example.com'))
b = bytes(dns_req)
p = DNS(b)
p.show()
```

Example 2:

```
msg =
b'\x00\x00\x01\x00\x00\x01\x00\x00\x00\x00\x00\x00\x03abc\x07example\x03net\x00\x00\x01\x00\x01'
DNS(msg).show()
```

2) Basic example on issuing HTTPS requests using the Requests library:

```
import requests
r = requests.get('https://1.1.1.1/')
r.status_code
```

You should get `r.status_code = 200`.

If the above works, then all the information you need to correctly make DoH requests can be found in these two links:

<https://developers.cloudflare.com/1.1.1.1/encrypted-dns/dns-over-https/make-api-requests/dns-wireformat>

<https://docs.python-requests.org/en/master/user/quickstart/>

NOTE: when working from the UGA campus network (including from your VM), use IP addresses (e.g., 1.1.1.1, 8.8.8.8, etc.) in the DoH query URL, because it appears that EITS is blocking some domain name strings related to DoH servers (e.g., dns.goole or cloudflare-dns.com).